

Trace: - application\_base

realtime:documentation:howto:applications:application\_base

## HOWTO build a simple RT application

The POSIX [API](#) forms the basis of real-time applications running under PREEMPT\_RT. For the real-time thread a POSIX thread is used (pthread). Every real-time application needs proper handling in several basic areas like scheduling, priority, memory locking and stack prefaulting.

### Basic prerequisites

Three basic prerequisites are introduced in the next subsections, followed by a short example illustrating those aspects.

### Scheduling and priority

The [scheduling policy](#) as well as the priority must be set by the application explicitly. There are two possibilities for this:

#### 1. Using `sched_setscheduler()`

This function needs to be called in the start routine of the pthread before calculating RT specific stuff.

#### 2. Using pthread attributes

The functions `pthread_attr_setschedpolicy()` and `pthread_attr_setschedparam()` offer the interfaces to set policy and priority. Furthermore scheduler inheritance needs to be set properly to `PTHREAD_EXPLICIT_SCHED` by using `pthread_attr_setinheritsched()`. This forces the new thread to use the policy and priority specified by the pthread attributes and not to use the inherit scheduling of the thread which created the real-time thread.

#### 3. Problems with pthread condition variables

Multithreaded applications which rely on glibc's libpthread are prone to unexpected latency delays since its condition variable implementation does not honor priority inheritance (🐞 [bugzilla](#)). Unfortunately glibc's [DNS](#) resolver and asynchronous I/O implementations depend in turn on these condition variables. 🐞 [librtpi](#) is an alternative [LGPL](#)-licensed pthread implementation which supports priority inheritance, and whose [API](#) is as close to glibc's as possible. The alternative 🐞 [MUSL libc](#) has a pthread condition variable implementation similar to glibc's.

### Memory locking

See [here](#)

### Stack for RT thread

See [here](#)

### Capabilities: running the app with RT priority as a non-root user

Several of the Pthread APIs, like `mlockall()`, `pthread_attr_setschedpolicy()`, by default and convention require root in order to successfully get their work done. Thus, RT apps - which need to set an RT sched policy and priority - are often run via `sudo`.

There's a far better approach to this; `sudo` gives the process root capabilities. This interests hackers 🤪. Instead, you should leverage the powerful POSIX **Capabilities** model! This way, the process (and threads) get `_only_` the capabilities they require and nothing more. This follows the infosec best practice, the *principle of least privilege*.

Apps start out with no capabilities by default; also note that capabilities are a per-thread resource (essentially translating to bitmasks with the task structure, which is per-thread of course). Among the various capability bits, the man page on `capabilities(7)` shows that `CAP_SYS_NICE` is the appropriate capability to use in this circumstance; a snippet from the `capabilities(7)` man page reveals this:

... **CAP\_SYS\_NICE**

- Lower the process nice value (`nice(2)`, `setpriority(2)`) and change the nice value for arbitrary processes;
- **set real-time scheduling policies for calling process, and set scheduling policies and priorities for arbitrary processes** (`sched_setscheduler(2)`, `sched_setparam(2)`, `sched_setattr(2)`);
- set CPU affinity for arbitrary processes (`sched_setaffinity(2)`);
- set I/O scheduling class and priority for arbitrary processes (`ioprio_set(2)`);
- apply `migrate_pages(2)` to arbitrary processes and allow processes to be migrated to arbitrary nodes;
- apply `move_pages(2)` to arbitrary processes;
- use the `MPOL_MF_MOVE_ALL` flag with `mbind(2)` and `move_pages(2)`. ...

**Ok, great, but how exactly is this capability bit to be set on the app?**

1. One approach is to do so programmatically, via the `capget()/capset()` system calls. (Note that it's generally easier to use the libcap library wrappers, `cap_[g|s]et_proc(3)`: 🐞 [https://man7.org/linux/man-pages/man3/cap\\_get\\_proc.3.html](https://man7.org/linux/man-pages/man3/cap_get_proc.3.html). This man page even provides a small example of doing so).
2. Another easy way is to leverage systemd and run your app as a service; in the service unit, specify the capability (see the man page on `systemd.exec(5)`: 🐞 <https://www.freedesktop.org/software/systemd/man/systemd.exec.html#Capabilities>).
3. Perhaps the easiest way: via the `setcap(8)` utility (it's man page: 🐞 <https://man7.org/linux/man-pages/man8/setcap.8.html>). The `setcap/getcap` are typically part of the libcap package. For example:

```
sudo setcap CAP_SYS_NICE+eip <your-app-binary-executable>
```

You could put this line in the app Makefile (or equivalent). (The `getcap(8)` utility can be used to verify that the 'dumb-capability' binary now has the `CAP_SYS_NICE` bit set)

And you're all set to run it as non-root now, a much more secure approach.

### Example

```
/*  
 * POSIX Real Time Example
```

#### Table of Contents

- HOWTO build a simple RT application
- Basic prerequisites
  - Scheduling and priority
  - Memory locking
  - Stack for RT thread
- Capabilities: running the app with RT priority as a non-root user
- Example

```

/* using a single pthread as RT thread
*/

#include <limits.h>
#include <pthread.h>
#include <sched.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>

void *thread_func(void *data)
{
    /* Do RT specific stuff here */
    return NULL;
}

int main(int argc, char* argv[])
{
    struct sched_param param;
    pthread_attr_t attr;
    pthread_t thread;
    int ret;

    /* Lock memory */
    if(mlockall(MCL_CURRENT|MCL_FUTURE) == -1) {
        printf("mlockall failed: %m\n");
        exit(-2);
    }

    /* Initialize pthread attributes (default values) */
    ret = pthread_attr_init(&attr);
    if (ret) {
        printf("init pthread attributes failed\n");
        goto out;
    }

    /* Set a specific stack size */
    ret = pthread_attr_setstacksize(&attr, PTHREAD_STACK_MIN);
    if (ret) {
        printf("pthread setstacksize failed\n");
        goto out;
    }

    /* Set scheduler policy and priority of pthread */
    ret = pthread_attr_setschedpolicy(&attr, SCHED_FIFO);
    if (ret) {
        printf("pthread setschedpolicy failed\n");
        goto out;
    }

    param.sched_priority = 80;
    ret = pthread_attr_setschedparam(&attr, &param);
    if (ret) {
        printf("pthread setschedparam failed\n");
        goto out;
    }

    /* Use scheduling parameters of attr */
    ret = pthread_attr_setinheritsched(&attr, PTHREAD_EXPLICIT_SCHED);
    if (ret) {
        printf("pthread setinheritsched failed\n");
        goto out;
    }


    /* Create a pthread with specified attributes */
    ret = pthread_create(&thread, &attr, thread_func, NULL);
    if (ret) {
        printf("create pthread failed\n");
        goto out;
    }

    /* Join the thread and wait until it is done */
    ret = pthread_join(thread, NULL);
    if (ret)
        printf("join pthread failed: %m\n");

out:
    return ret;
}

```

realtime/documentation/howto/applications/application\_base.txt · Last modified: 2024/05/31 15:54 by alison

Except where otherwise noted, content on this wiki is licensed under the following license:  CC Attribution-Noncommercial-Share Alike 4.0 International